



November 5, 2018

Managed ML Model Training and Serving

Lisa Quera, ML Specialist

Google Cloud

Machine Learning at Scale



Cloud ML Engine



Serverless, no-ops, ML training and serving platform



Distributed training infrastructure that supports CPUs, GPUs, and TPUs



Automatic hyperparameter tuning



Train, tune, and serve TensorFlow models (batch and online prediction)



Train and serve Scikit-learn and XGBoost models for online predictions

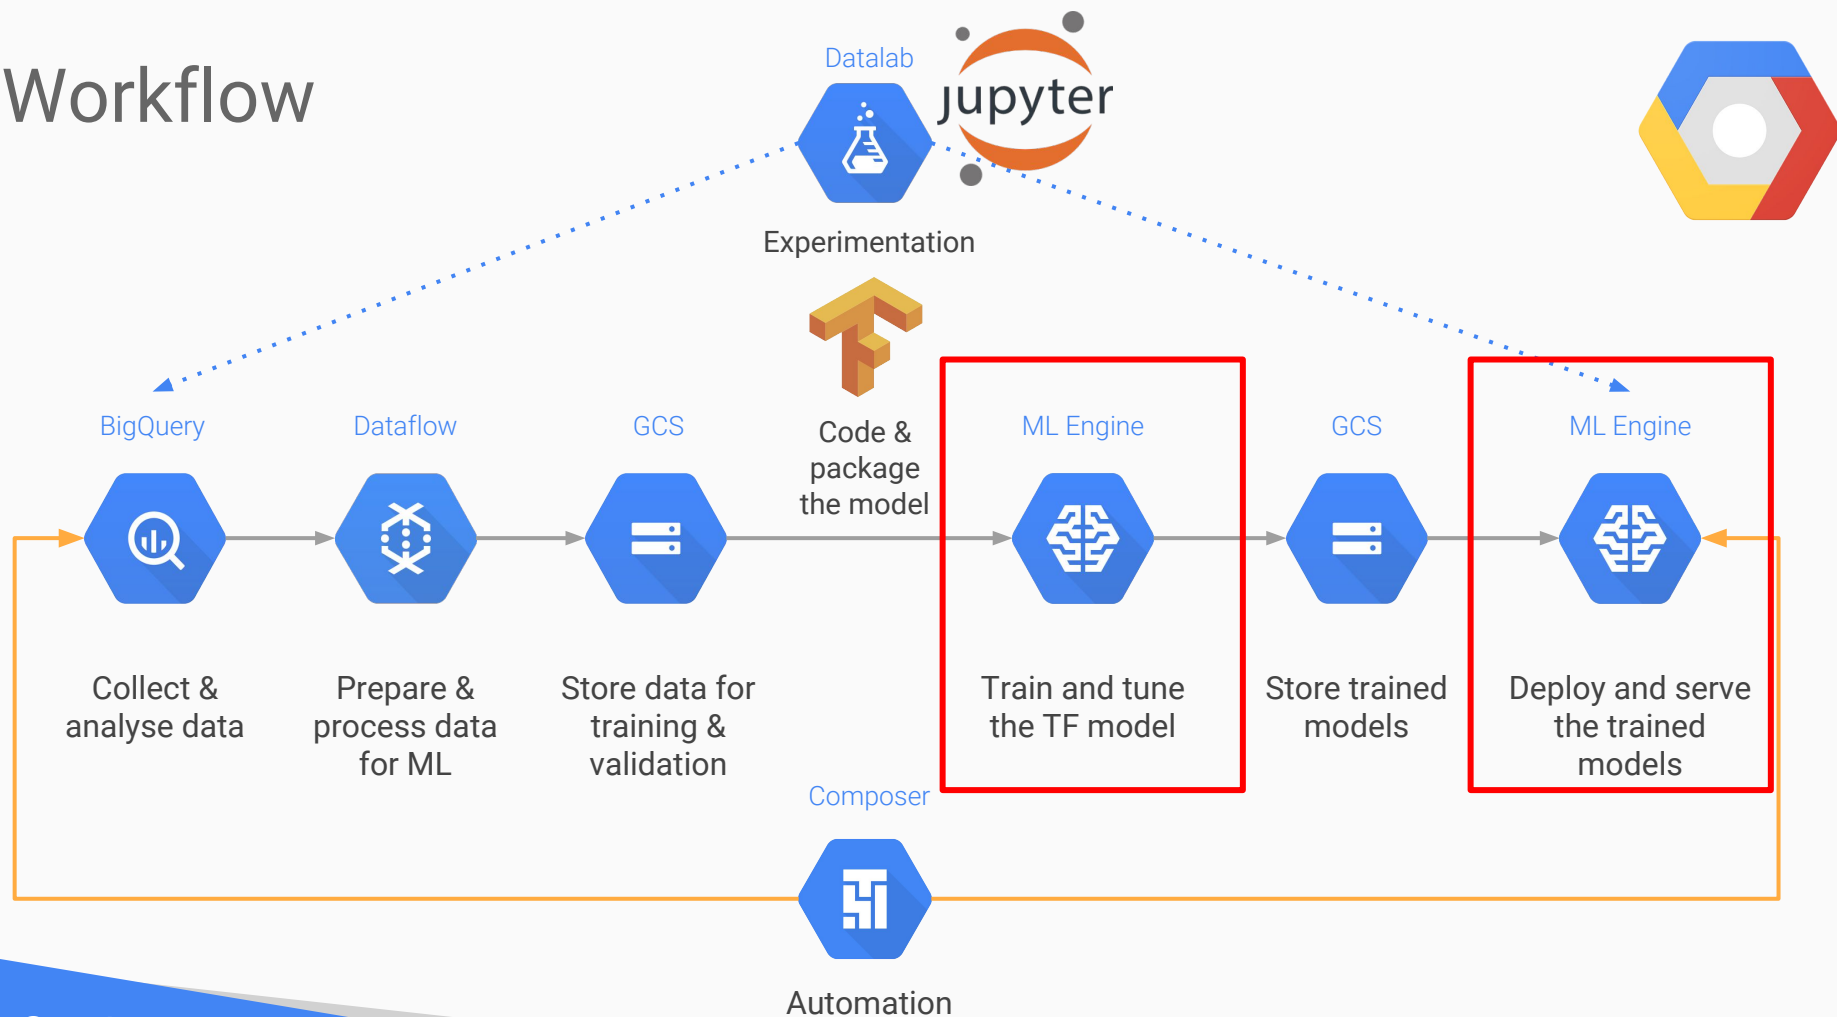


TPUs for training TensorFlow models **(Beta)**

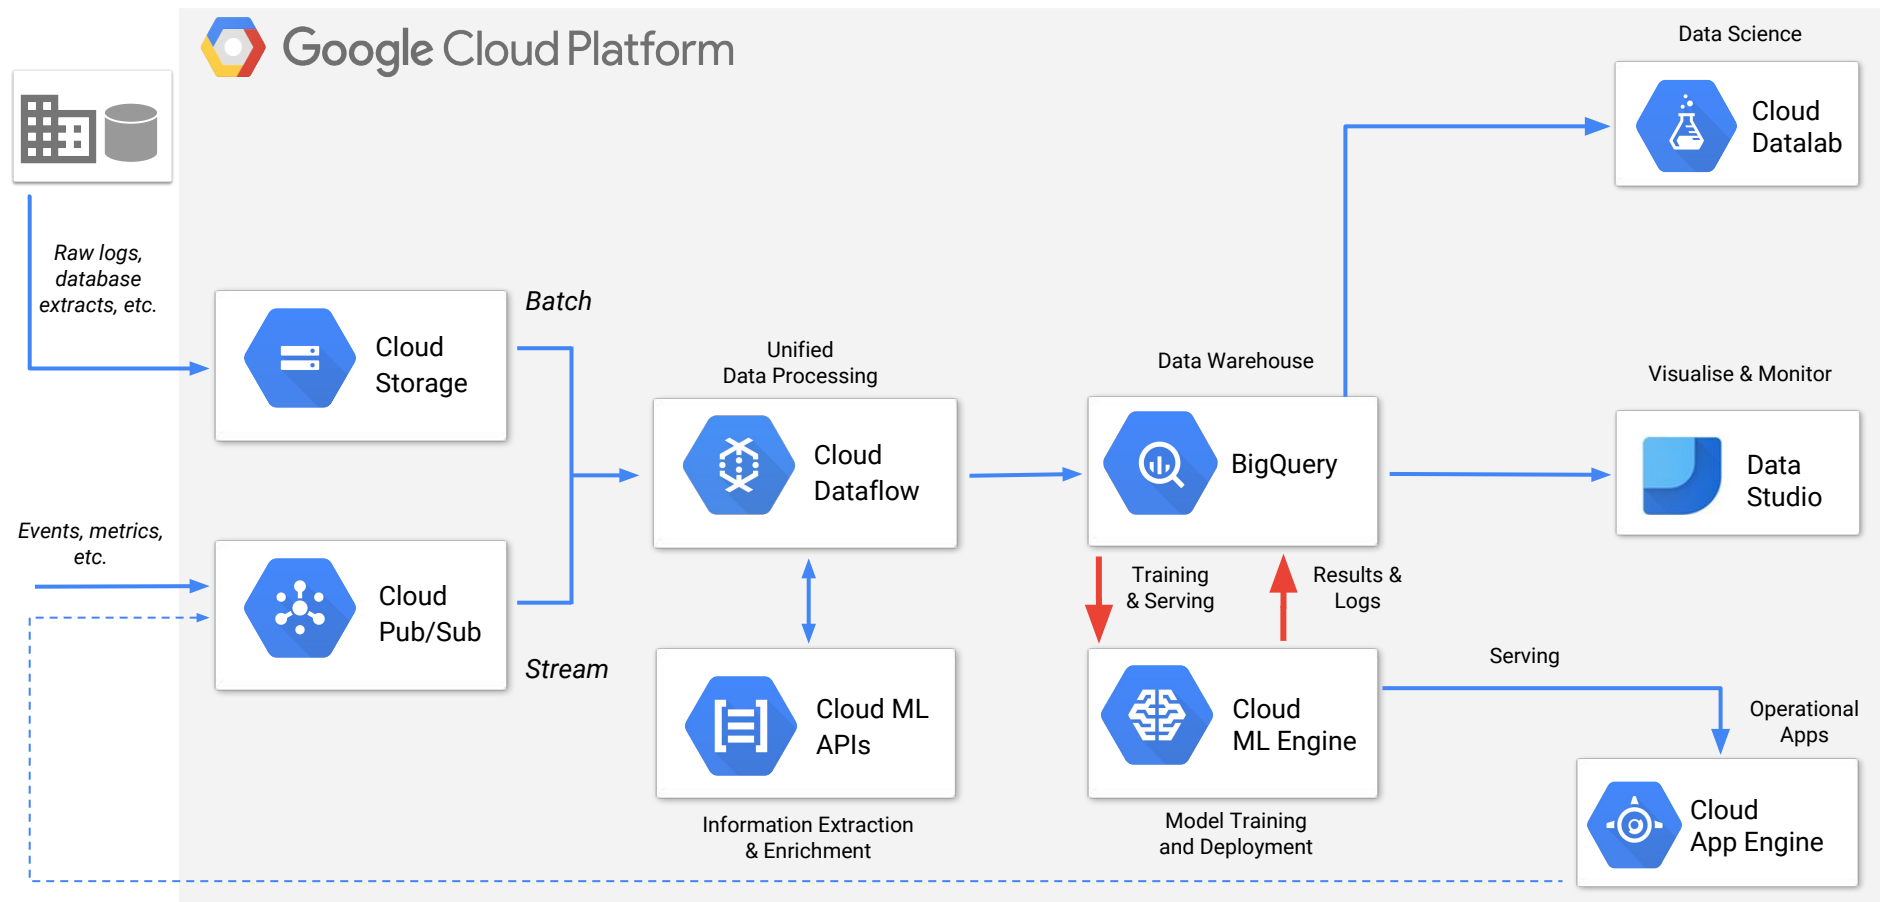


Google Cloud

Workflow



Example Architecture: BigQuery to deployed model API



Training locally



train locally

```
gcloud ml-engine local train \  
  --module-name trainer.task --package-path trainer/ \  
  -- \  
  --train-files $TRAIN_DATA --eval-files $EVAL_DATA --job-dir $MODEL_DIR
```

Local path

output directory

training data

evaluation data



Training in the cloud

with single node



train in the cloud

GCS location

Can be a package in GCS

region

```
gcloud ml-engine jobs submit training $JOB_NAME --job-dir $OUTPUT_PATH \
  --runtime-version 1.10 --module-name trainer.task --package-path trainer --region $REGION \
  --scale-tier BASIC --train-files $TRAIN_DATA --eval-files $EVAL_DATA --num-epoch 1000 --learning-rate 0.01
```

single worker

GCS locations

model-specific params



Google Cloud

<https://cloud.google.com/ml-engine/docs/tensorflow/machine-types>



Training in the cloud at scale

with multiple workers



<https://cloud.google.com/ml-engine/docs/tensorflow/machine-types>

```
gcloud ml-engine jobs submit training $JOB_NAME --job-dir $OUTPUT_PATH \  
  --runtime-version 1.10 --module-name trainer.task --package-path trainer --region $REGION \  
  --scale-tier STANDARD_1  
  -- \
```

distributed

```
--train-files $TRAIN_DATA --eval-files $EVAL_DATA
```



Manually Distributing the Training

```
with tf.device("/job:ps/task:0"):
    weights_1 = tf.Variable(...)
    biases_1 = tf.Variable(...)

with tf.device("/job:ps/task:1"):
    weights_2 = tf.Variable(...)
    biases_2 = tf.Variable(...)

with tf.device("/job:worker/task:7"):
    input, labels = ...
    layer_1 = tf.nn.relu(tf.matmul(input, weights_1) + biases_1)
    logits = tf.nn.relu(tf.matmul(layer_1, weights_2) + biases_2)
    # ...
    train_op = ...

with tf.Session("grpc://worker7.example.com:2222") as sess:
    for _ in range(10000):
        sess.run(train_op)
```

tf.train.ClusterSpec construction

```
tf.train.ClusterSpec({"local": ["localhost:2222", "localhost:2223"]})
```

```
tf.train.ClusterSpec({
    "worker": [
        "worker0.example.com:2222",
        "worker1.example.com:2222",
        "worker2.example.com:2222"
    ],
    "ps": [
        "ps0.example.com:2222",
        "ps1.example.com:2222"
    ]
})
```


Training in the cloud at scale

with GPUs (K80/P100/V100 - *availability by region*)



<https://cloud.google.com/ml-engine/docs/tensorflow/machine-types>

```
gcloud ml-engine jobs submit training $JOB_NAME --job-dir $OUTPUT_PATH \  
  --runtime-version 1.10 --module-name trainer.task --package-path trainer --region $REGION \  
  --scale-tier BASIC_GPU  
  -- \  
  --train-files $TRAIN_DATA --eval-files $EVAL_DATA
```

single GPU



CMLE Machine Types

Cloud ML Engine scale tier	
BASIC	<p>A single worker instance. This tier is suitable for learning how to use Cloud ML Engine and for experimenting with new models using small datasets.</p> <p>Compute Engine machine name: n1-standard-4</p>
STANDARD_1	<p>One master instance, plus four workers and three parameter servers.</p> <p>Compute Engine machine name, master: n1-highcpu-8, workers: n1-highcpu-8, parameter servers: n1-standard-4</p>
PREMIUM_1	<p>One master instance, plus 19 workers and 11 parameter servers.</p> <p>Compute Engine machine name, master: n1-highcpu-16, workers: n1-highcpu-16, parameter servers: n1-highmem-8</p>
BASIC_GPU	<p>A single worker instance with a single NVIDIA Tesla K80 GPU. To learn more about graphics processing units (GPUs), see the section on training with GPUs.</p> <p>Compute Engine machine name: n1-standard-8 with one k80 GPU</p>
BASIC_TPU <i>(Beta)</i>	<p>A master VM and a Cloud TPU. See how to use TPUs for your training job.</p> <p>Compute Engine machine name, master: n1-standard-4, workers: Cloud TPU</p>
CUSTOM	<p>The CUSTOM tier is not a set tier, but rather enables you to use your own cluster specification. When you use this tier, set values to configure your processing cluster according to these guidelines:</p>

CMLE Custom Options



Cloud ML Engine machine name	
standard	<p>A basic machine configuration suitable for training simple models with small to moderate datasets.</p> <p>Compute Engine machine name: n1-standard-4</p>
large_model	<p>A machine with a lot of memory, specially suited for parameter servers when your model is large (having many hidden layers or layers with very large numbers of nodes).</p> <p>Compute Engine machine name: n1-highmem-8</p>
complex_model_s	<p>A machine suitable for the master and workers of the cluster when your model requires more computation than the standard machine can handle satisfactorily.</p> <p>Compute Engine machine name: n1-highcpu-8</p>
complex_model_m	<p>A machine with roughly twice the number of cores and roughly double the memory of complex_model_s.</p> <p>Compute Engine machine name: n1-highcpu-16</p>
complex_model_l	<p>A machine with roughly twice the number of cores and roughly double the memory of complex_model_m.</p> <p>Compute Engine machine name: n1-highcpu-32</p>
standard_gpu	<p>A machine equivalent to standard that also includes a single NVIDIA Tesla K80 GPU.</p> <p>Compute Engine machine name: n1-standard-8 with one k80 GPU</p>
complex_model_m_gpu	<p>A machine equivalent to complex_model_m that also includes four NVIDIA Tesla K80 GPUs.</p> <p>Compute Engine machine name: n1-standard-16-k80x4</p>

Cloud TPUs



Cloud TPUs

- \$ Excellent performance / \$
- 🕒 Train in days instead of weeks
- ✓ No more fighting with drivers
- 📄 Flexibility and scale
- ☑ Fully-managed in the cloud

Supported models for TPUs



Image recognition & object detection

Image recognition:

AmoebaNet-D
ResNet-50/101/152/200
Inception v2/v3/v4
DenseNet

Object detection:

RetinaNet

Low-resource models:

MobileNet
SqueezeNet



Machine translation and language modeling

Models:

Machine translation
Language modeling
Sentiment analysis
Question-answering
(all transformer-based)



Speech recognition

Model:

ASR Transformer
(LibriSpeech)



Image generation

Models:

Image Transformer
DCGAN

Training Scikit-learn & XGBoost models




*Key parameters to scikit-learn
and XGBoost training on CMLE*

```
gcloud ml-engine jobs submit training $JOB_NAME --job-dir $OUTPUT_PATH \  
  --runtime-version 1.9 --python-version 2.7 --scale-tier BASIC \  
  --module-name sklearn_trainer.task --package-path sklearn_trainer --region $REGION \  
  -- \  
  --train-files $TRAIN_DATA --eval-files $EVAL_DATA --num-epoch 1000 --learning-rate 0.01
```

Keras: what is happening?



- Compatibility module introduced in TensorFlow: `tf.keras`
- Write your custom estimator `model_fn` using `tf.keras.layers` and/or `tf.layers` (mix-and-match)
- Convert your compiled keras model to `tf.estimator` using `tf.keras.estimator.model_to_estimator`
- With TensorFlow, Keras users gain access to new features:
 - Distributed training
 - Multiple GPUs
 - Cloud ML
 - Hyperparameter tuning
 - TF-Serving



Keras

Keras example




```
y_train = ...
y_test = ...

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))


model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(), metrics=['accuracy'])

model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
```



Keras

Premade Estimators



TensorFlow



LinearClassifier



DNNClassifier



DNNLinearCombinedClassifier



DNNLinearCombinedRegressor



(tf.contrib.kernel_methods.) KernelLinearClassifier



(tf.contrib.factorization.) KMeansClustering



(tf.contrib.timeseries.) ARRegressor



BoostedTreesClassifier



BoostedTreesRegressor



LinearRegressor



DNNRegressor



(tf.contrib.estimator)RNNClassifier



(tf.contrib.estimator)RNNEstimator

Hyperparameter tuning



```
gcloud ml-engine jobs submit training $JOB_NAME --job-dir $OUTPUT_PATH \  
  --runtime-version 1.7 --module-name trainer.task --package-path trainer/ --region $REGION \  
  --scale-tier PREMIUM_1 --config hyperparams.yaml  
  -- \  
  --train-files $TRAIN_DATA --eval-files $EVAL_DATA
```

hypertuning



Hyperparameter tuning



hyperparams.yaml

```
trainingInput:
  hyperparameters:
    goal: MAXIMIZE
    hyperparameterMetricTag: accuracy
    maxTrials: 40
    enableTrialEarlyStopping: True
    maxParallelTrials: 2
    algorithm: UNSPECIFIED
    params:
      - parameterName: learning-rate
        type: FLOAT
        minValue: 0.001
        maxValue: 0.1
        scaleType: UNIT_LOG_SCALE
```

...



Google Cloud

task.py

```
...

# Initialise the optimizer for the DNN
optimizer = tf.train.AdagradOptimizer(
    learning_rate=hparams.learning_rate)

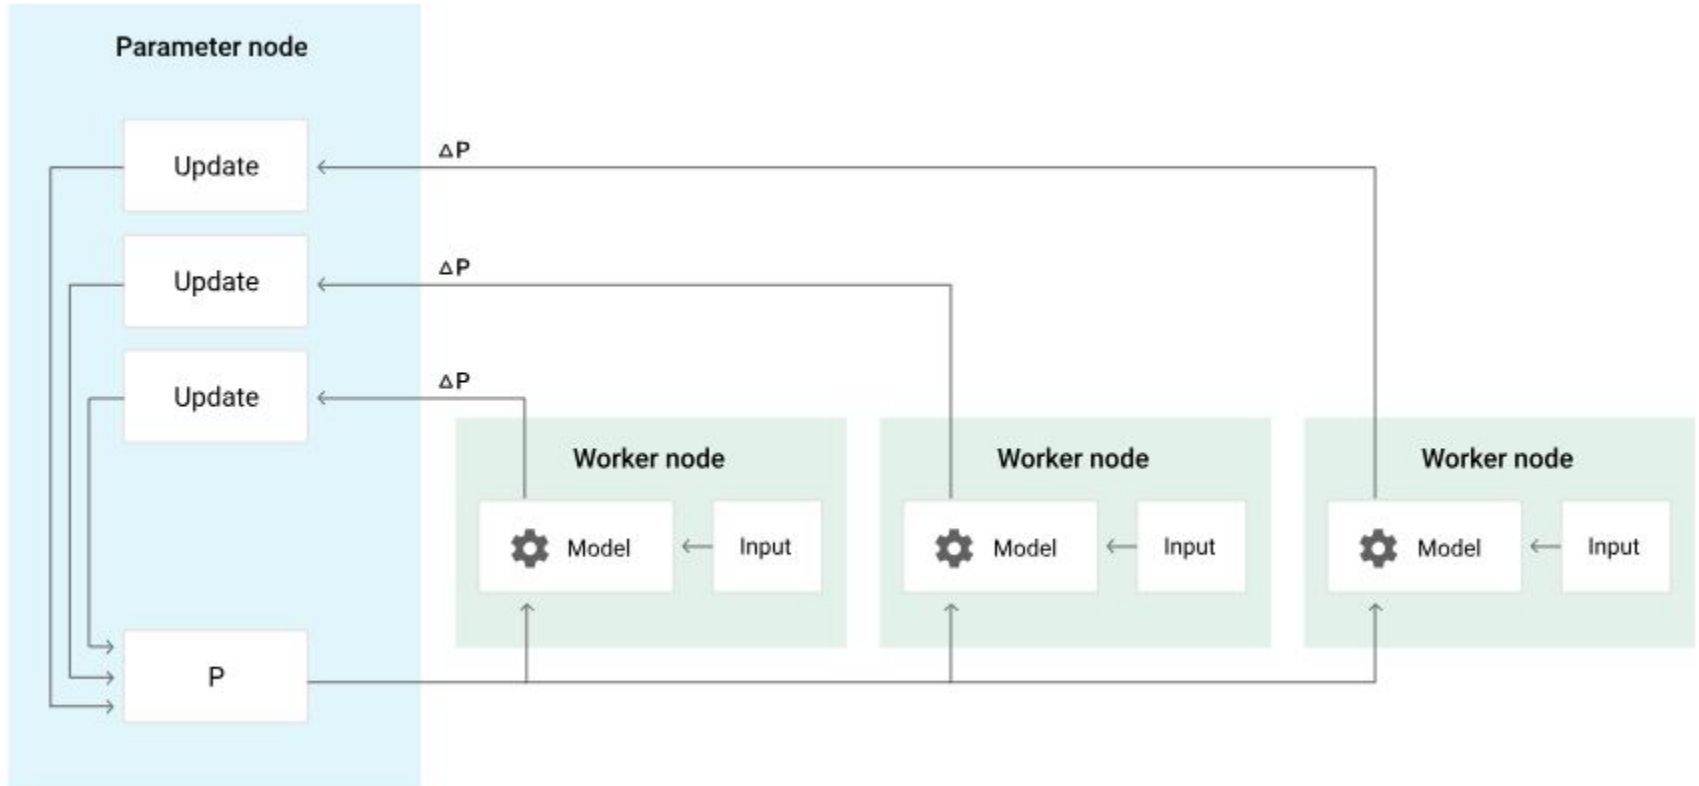
...

parser.add_argument(
    '--learning-rate',
    help='Learning rate used by the DNN optimizer',
    default=0.01,
    type=float
)

...
```



Distributed Training





Consuming the deployed
ML model API for
predictions

Deploy the trained TF model



gcloud command line tool:

```
# Creating model
NAME=demo_classifier
gcloud ml-engine models create $NAME --regions $REGION

# Creating versions\
VERSION=v2.3
MODEL_DIR=gs://ksalama-gcs/trained_models/demo_classifier_output

gcloud ml-engine versions create $VERSION --model $NAME --origin $MODEL_DIR \
  --runtime-version 1.7 --config config.yaml
```

```
# List deployed models
gcloud ml-engine models list
```



description: A free-form description of the version.
deploymentUri: gs://path/to/source
runtimeVersion: '1.7'
manualScaling:
 nodes: 10
autoScaling:
 minNodes: 0



Predicting with TF Model



Online versus Batch Prediction

[SEND FEEDBACK](#)

Internal: Count: 1, Average: 5.0

Cloud ML Engine provides two ways to get predictions from trained models: *online prediction* (sometimes called HTTP prediction), and *batch prediction*. In both cases, you pass input data to a cloud-hosted machine-learning model and get inferences for each data instance. The differences are shown in the following table:

Online prediction	Batch prediction
Optimized to minimize the latency of serving predictions.	Optimized to handle a high volume of instances in a job and to run more complex models.
Can process one or more instances per request.	Can process one or more instances per request.
Predictions returned in the response message.	Predictions written to output files in a Cloud Storage location that you specify.
Input data passed directly as a JSON string.	Input data passed indirectly as one or more URIs of files in Cloud Storage locations.
Returns as soon as possible.	Asynchronous request.
Accounts with the following IAM roles can request online predictions: <ul style="list-style-type: none">• Legacy Editor or Viewer• Cloud ML Engine Admin or Developer	Accounts with the following IAM roles can request batch predictions: <ul style="list-style-type: none">• Legacy Editor• Cloud ML Engine Admin or Developer
Runs on the runtime version and in the region selected when you deploy the model.	Can run in any available region, using any available runtime version. Though you should run with the defaults for deployed model versions.

Predicting with TF Model



gcloud command line tool:

```
gcloud ml-engine predict --model $NAME --version $VERSION --json-instances test.json
```

gcloud **batch prediction**:

```
gcloud ml-engine job submit prediction \  
$JOB_NAME --model $NAME --version $VERSION \  
data-format TEXT \  
input-paths $GCS_DATA_DIR \  
output-path $GCS_OUT_DIR \  

```

python code - REST API call - **online prediction**:

```
def estimate(project, model_name, version, instances):  
    credentials = GoogleCredentials.get_application_default()  
    api = discovery.build('ml', 'v1',  
                          credentials=credentials,  
                          discoveryServiceUrl =  
                          'https://storage.googleapis.com/cloud-ml/discovery/ml_v1_discovery.json')  
  
    request_data = {'instances': instances}  
  
    model_url = 'projects/{}/models/{}/versions/{}'.format(  
                project, model_name, version)  
  
    response = api.projects().predict(body=request_data, name=model_url).execute()  
  
    estimates = list(map(lambda item: item["scores"],  
                        response["predictions"]  
                        ))  
  
    return estimates
```



Deploy XGBoost & Scikit-learn models



```
from sklearn.externals import joblib
joblib.dump(estimator, model.joblib)
```

*Save and dump
model to GCS*

```
gsutil cp ./model.joblib ${MODEL_PATH}/model.joblib
```

*deploy the model
to CMLE*

```
gcloud ml-engine models create ${MODEL_NAME} --regions=${REGION}
```




```
gcloud ml-engine versions create ${VERSION} --model=${MODEL_NAME} \
--origin=${MODEL_PATH} \
--runtime-version="1.4" \
--framework="SCIKIT_LEARN"
--pythonVersion="2.7"
```

<https://cloud.google.com/blog/big-data/2018/04/serving-real-time-scikit-learn-and-xgboost-predictions>



Monitoring model API serving health



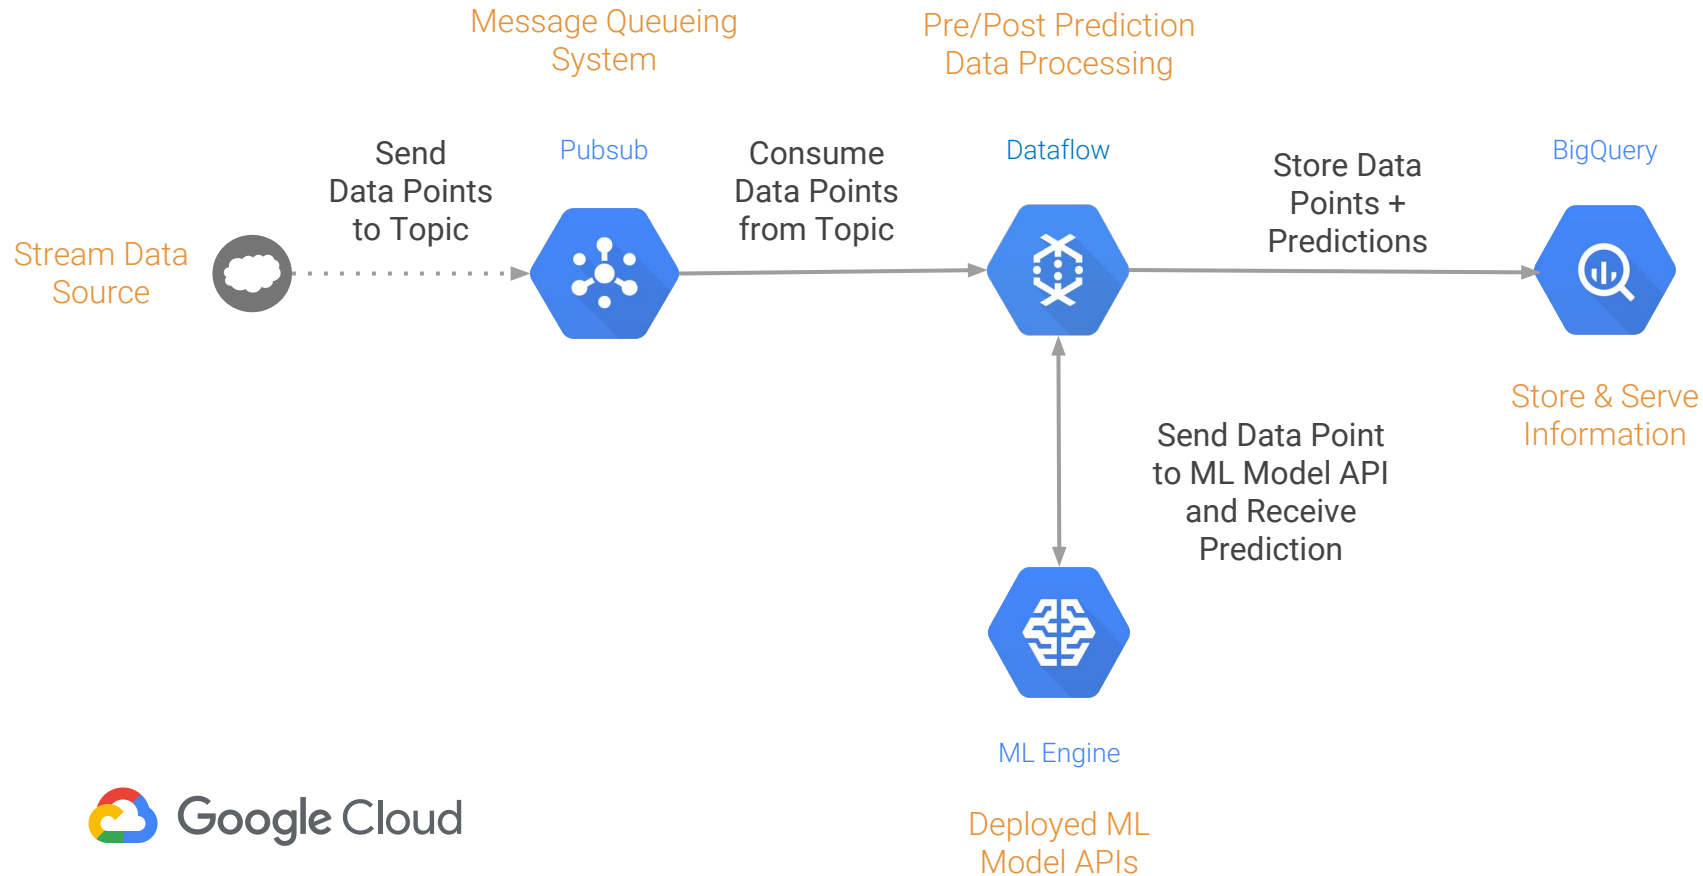
✓ v1



Description	
Model	taxifare_estimator
Model location	gs://ksalama-gcs-cloudml/ml-models/taxifare/dnn-combined-regression-small,
Creation time	Aug 22, 2017, 11:52:04 PM
Last use time	Oct 11, 2017, 1:20:01 PM



Example Production Workflow



Monitor ML Engine job on Cloud Console

Stackdriver
Logging

Logs

Logs-based metrics

Exports

Logs ingestion

CREATE METRIC

CREATE EXPORT

↺

▶

Filter by label or text search

Cloud ML Job, lab3a_181009_204525

All logs

Any log level

No limit

Jump to now

Showing logs from all time (PST)

Download logs

View Options

▶	2018-10-09 14:11:42.628 PDT	master-replica-0	Finished evaluation at 2018-10-09-21:11:42	⋮
▶	2018-10-09 14:11:42.629 PDT	master-replica-0	Saving dict for global step 96544: average_loss = 89.4802, gl...	⋮
▶	2018-10-09 14:11:43.108 PDT	master-replica-0	Restoring parameters from gs://lquera-workshop-gcs/taxifare/c...	⋮
▶	2018-10-09 14:11:43.796 PDT	master-replica-0	Assets added to graph.	⋮
▶	2018-10-09 14:11:43.796 PDT	master-replica-0	No assets to write.	⋮
▶	2018-10-09 14:11:47.008 PDT	master-replica-0	SavedModel written to: gs://lquera-workshop-gcs/taxifare/ch3/...	⋮
▶	2018-10-09 14:11:48.288 PDT	master-replica-0	global_step/sec: 30.1842	⋮
▶	2018-10-09 14:11:48.293 PDT	master-replica-0	Saving checkpoints for 100007 into gs://lquera-workshop-gcs/t...	⋮
▶	2018-10-09 14:11:52.836 PDT	master-replica-0	Skip the current checkpoint eval due to throttle secs (300 se...	⋮
▶	2018-10-09 14:11:53.177 PDT	master-replica-0	Loss for final step: 40353.3.	⋮
▶	2018-10-09 14:11:53.177 PDT	master-replica-0	Training has already ended. But the last eval is skipped due ...	⋮
▶	2018-10-09 14:11:53.743 PDT	master-replica-0	Starting evaluation at 2018-10-09-21:11:53	⋮
▶	2018-10-09 14:11:53.803 PDT	master-replica-0	Restoring parameters from gs://lquera-workshop-gcs/taxifare/c...	⋮
▶	2018-10-09 14:13:40.509 PDT	master-replica-0	Finished evaluation at 2018-10-09-21:13:40	⋮
▶	2018-10-09 14:13:40.509 PDT	master-replica-0	Saving dict for global step 100007: average_loss = 89.4804, g...	⋮
▶	2018-10-09 14:13:41.016 PDT	master-replica-0	Restoring parameters from gs://lquera-workshop-gcs/taxifare/c...	⋮
▶	2018-10-09 14:13:41.666 PDT	master-replica-0	Assets added to graph.	⋮
▶	2018-10-09 14:13:41.667 PDT	master-replica-0	No assets to write.	⋮
▶	2018-10-09 14:13:44.554 PDT	master-replica-0	SavedModel written to: gs://lquera-workshop-gcs/taxifare/ch3/...	⋮
▶	2018-10-09 14:13:45.848 PDT	master-replica-0	Module completed; cleaning up.	⋮
▶	2018-10-09 14:13:45.848 PDT	master-replica-0	Clean up finished.	⋮
▶	2018-10-09 14:13:45.849 PDT	master-replica-0	Task completed successfully.	⋮
▶	2018-10-09 14:14:00.044 PDT	master-replica-0	Tearing down training program.	⋮

Monitor ML Engine job on Cloud Console

Stackdriver
Logging

Logs

Logs-based metrics

Exports

Logs ingestion

CREATE METRIC

CREATE EXPORT

text:average_loss

Cloud ML Job, lab3a_181009_204525

All logs

Any log level

No limit

Jump to now

Showing logs from all time (PST)

Download logs

View Options

No older entries found matching current filter.

2018-10-09 13:51:40.778 PDT

master-replica-0

Saving dict for global step 18: average_loss = 212.268, globa...

2018-10-09 14:01:39.560 PDT

master-replica-0

Saving dict for global step 47742: average_loss = 89.53, glob...

2018-10-09 14:11:42.629 PDT

master-replica-0

Saving dict for global step 96544: average_loss = 89.4802, gl...

2018-10-09 14:13:40.509 PDT

master-replica-0

Saving dict for global step 100007: average_loss = 89.4804, g...

Load newer logs

Thank you!

